

Homework Assignment 9

Computer Vision for HCI

Prof. Jim Davis

TA: Sayan Mandal

Anirudh Ganesh

CSE5524 (Au '18)

Score: ___/12

Due Date: 10/30/18

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from skimage.io import imread
from scipy.signal import fftconvolve
from sklearn.neighbors import KNeighborsClassifier
from matplotlib.colors import ListedColormap
```

1 Part 1 Template Matching NCC

```
In [2]: search = imread('./data/search.png')
template = imread('./data/template.png')
```

```
In [3]: image_shape = search.shape
```

```
image = np.array(search, dtype=np.float64, copy=False)
```

```
In [4]: pad_width = tuple((width, width) for width in template.shape)
```

```
In [5]: image = np.pad(image, pad_width=pad_width, mode='constant', constant_values=0)
```

```
In [6]: def calc_window_sum(image, window_shape):
    window_sum = np.cumsum(image, axis=0)
    window_sum = (window_sum[window_shape[0]:-1] - window_sum[:-window_shape[0] - 1])
    window_sum = np.cumsum(window_sum, axis=1)
    window_sum = (window_sum[:, window_shape[1]:-1] - window_sum[:, :-window_shape[1] - 1])
    window_sum = np.cumsum(window_sum, axis=2)
    window_sum = (window_sum[:, :, window_shape[2]:-1] - window_sum[:, :, :-window_shape[2] - 1])
    return window_sum
```

```
In [7]: image_window_sum = calc_window_sum(image, template.shape)
image_window_sum2 = calc_window_sum(image ** 2, template.shape)
```

```
In [8]: template_mean = template.mean()
template_volume = np.prod(template.shape)
template_ssd = np.sum((template - template_mean) ** 2)
```

```
In [9]: xcorr = fftconvolve(image, template[::-1, ::-1, ::-1], mode="valid")[1:-1, 1:-1, 1:-1]
```

```
In [10]: numerator = xcorr - image_window_sum * template_mean
```

```
In [11]: denominator = image_window_sum2
```

```
In [12]: image_window_sum = np.multiply(image_window_sum, image_window_sum)
```

```
In [13]: image_window_sum = np.divide(image_window_sum, template_volume)
```

```

In [14]: denominator -= image_window_sum
         denominator *= template_ssd

In [15]: denominator = np.maximum(denominator, 0) # sqrt of negative number not allowed

In [16]: denominator = np.sqrt(denominator)

In [17]: response = np.zeros_like(xcorr, dtype=np.float64)

In [18]: mask = denominator > np.finfo(np.float64).eps

In [19]: response[mask] = numerator[mask] / denominator[mask]

In [20]: slices = []

In [21]: for i in range(template.ndim):
         d0 = template.shape[i] - 1
         d1 = d0 + image_shape[i] - template.shape[i] + 1
         slices.append(slice(d0, d1))

In [22]: result = response[tuple(slices)]

In [23]: def largest_indices(ary, n):
         """Returns the n largest indices from a numpy array."""
         flat = ary.flatten()
         indices = np.argpartition(flat, -n)[-n:]
         indices = indices[np.argsort(-flat[indices])]
         return np.unravel_index(indices, ary.shape)

In [24]: ind_list = [1, 2, 5, 10, 100, 500]
         x_list = []
         y_list = []

         for t in ind_list:
             ind = largest_indices(result, t)
             x_t, y_t = ind[1][-1], ind[0][-1]
             x_list.append(x_t)
             y_list.append(y_t)

In [25]: ij = np.unravel_index(np.argmax(result), result.shape)
         _, x, y = ij[::-1] # IMPORTANT TO SWITCH ROWS AND COLUMNS I've LOST 5 points this seme

In [26]: plot_n = 6

         fig, axarr = plt.subplots(6, 2, figsize=(15,20))

         for i in range(plot_n):
             htemplate, wtemplate, _ = template.shape

             axarr[i, 0].imshow(search[y_list[i]: y_list[i]+htemplate, x_list[i]:x_list[i]+wtemp

```

```

axarr[i, 0].set_axis_off()
axarr[i, 0].set_title(f'detected region, {ind_list[i]} closest')

axarr[i, 1].imshow(search)
axarr[i, 1].set_axis_off()
axarr[i, 1].set_title(f'image: ({x_list[i]}, {y_list[i]}')')

rect = plt.Rectangle((x_list[i], y_list[i]), wtemplate, htemplate, edgecolor='r', f

axarr[i, 1].add_patch(rect)


plt.show()

```

detected region, 1 closest




image: (71,82)



detected region, 2 closest




image: (70,81)



detected region, 5 closest



image: (70,80)



detected region, 10 closest






image: (71,84)



detected region, 100 closest



image: (66,83)



detected region, 500 closest



image: (200,67)



We can see that our template matching performs really well in that it identifies “Stampy” the elephant pretty much pixel for pixel. Even the 2nd and 5th closest guesses are almost indistinguishable because of their 1-2 pixel discrepancy. We see however that in 100th closest our match drifts little too much to the left, as witnessed by the fact that there is more of John Swartzwelder (the guy in suit) and Brandine (the lady on the right) is now nowhere to be seen. By 500th closest, we can clearly see that the matching is way off, (its now closer to Homer rather than Stampy). This is in line with what we discussed in the slides with the person’s image and the mountain behind

2 Part 2 Template Matching Stereo Vision

```
In [27]: left = imread('./data/left.png')
         right = imread('./data/right.png')
```

```
In [28]: def stereo_match(left_img, right_img, kernel=11, max_offset=50):
```

```
    w, h = left_img.shape
    depth = np.zeros((w, h), np.uint8)
    depth.shape = h, w
    kernel_half = kernel // 2
    offset_adjust = 255 / max_offset
```

```
    for y in range(kernel_half, h - kernel_half):
```

```
        for x in range(kernel_half, w - kernel_half):
            best_offset = 0
            prev_ncc = float("-inf")
```

```
            for offset in range(max_offset):
```

```
                ncc = 0
```

```
                lwindow = left[y-kernel_half:y+kernel_half, x-kernel_half:x+kernel_half]
                rwindow = right[y-kernel_half:y+kernel_half, x-kernel_half:x+kernel_half]
```

```
                lwindow_mean = lwindow.mean()
                lwindow_std = np.std(lwindow)
                rwindow_mean = rwindow.mean()
                rwindow_std = np.std(rwindow)
```

```
                for v in range(-kernel_half, kernel_half):
```

```
                    for u in range(-kernel_half, kernel_half):
```

```
                        ncc_temp = (int(left[y+v, x+u]) - lwindow_mean) * (int(right[y+v,
```

```
                        if lwindow_std ==0 or rwindow_std ==0:
```

```
                            ncc_temp = 0
```

```
                        else:
```

```
ncc_temp /= lwindow_std*rwindow_std
ncc += ncc_temp
```

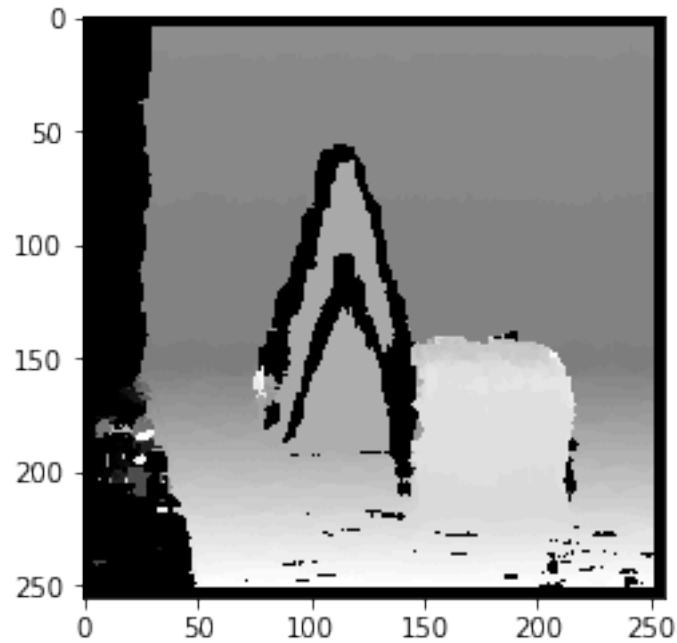
```
if ncc > prev_ncc:
    prev_ncc = ncc
    best_offset = offset
```

```
depth[y, x] = best_offset * offset_adjust
```

```
return depth
```

```
In [30]: plt.imshow(stereo_match(left, right), cmap='gray')
```

```
Out[30]: <matplotlib.image.AxesImage at 0x1cb88eaffd0>
```



3 Part 3 k-Nearest-Neighbours

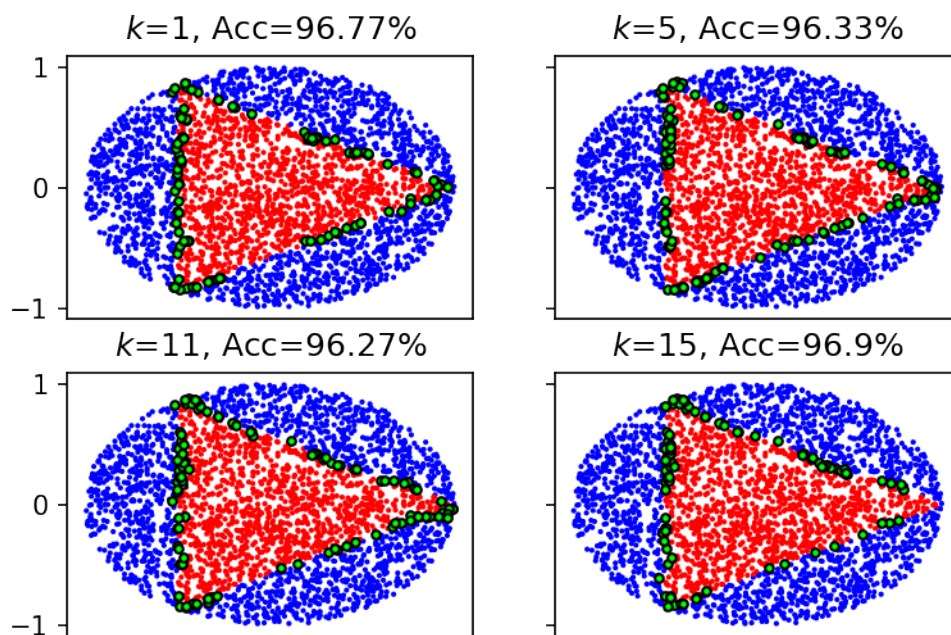
```
In [31]: train = np.loadtxt('./data/train.txt')
         test = np.loadtxt('./data/test.txt')
         X_train, y_train = train[:,0:2], train[:,2]
         X_test, y_test = test[:,0:2], test[:,2]
```

```
In [32]: neighbours = [1, 5, 11, 15]
```

```

In [33]: f, axarr = plt.subplots(2, 2, sharex='col', sharey='row', dpi=150)
         for n in neighbours:
             idx = neighbours.index(n)
             model = KNeighborsClassifier(n_neighbors=n)
             model.fit(X_train,y_train)
             y_pred = model.predict(X_test)
             cmap_bold = ListedColormap(['#FF0000', '#00FF00', '#0000FF'])
             axarr[idx//2,idx%2].tick_params(axis='x', which='both', bottom=False, top=False, labelsize=10)
             axarr[idx//2,idx%2].scatter(X_test[:, 0], X_test[:, 1], c=y_pred, cmap=cmap_bold, s=10)
             axarr[idx//2,idx%2].scatter(X_test[y_pred!=y_test, 0], X_test[y_pred!=y_test, 1], c='g', s=10)
             axarr[idx//2,idx%2].set_title(f"$k$={n}, Acc={round(model.score(X_test, y_test)*100)}%")
         plt.show()

```



We notice that most of the errors made in classification, it is at the border of our classes, this is expected as that is where the nearest neighbours will be the most contested. However as we increase our k , at first we notice that our accuracy decreases, as in this case there might be cases where more neighbours cause the classifier to be misled. But once we reach $k=15$, we see that our accuracy jumps up, thus indicating that at this stage we are getting sufficient enough information from our neighbours that its indicative of our overall sample.